

# Week 3 homework 1

Alex Parra

11/6/2022

```
library(tsibble)
```

```
## Warning: package 'tsibble' was built under R version 4.1.3
```

```
##
```

```
## Attaching package: 'tsibble'
```

```
## The following objects are masked from 'package:base':
```

```
##
```

```
## intersect, setdiff, union
```

```
library(fpp3)
```

```
## Warning: package 'fpp3' was built under R version 4.1.3
```

```
## -- Attaching packages ----- fpp3 0.4.0 --
```

```
## v tibble      3.1.6      v ggplot2      3.3.5
```

```
## v dplyr       1.0.7      v tsibbledata 0.4.0
```

```
## v tidyr       1.2.0      v feasts      0.2.2
```

```
## v lubridate   1.8.0      v fable       0.3.1
```

```
## Warning: package 'tsibbledata' was built under R version 4.1.3
```

```
## Warning: package 'feasts' was built under R version 4.1.3
```

```
## Warning: package 'fabletools' was built under R version 4.1.3
```

```
## Warning: package 'fable' was built under R version 4.1.3
```

```
## -- Conflicts ----- fpp3_conflicts --
```

```
## x lubridate::date()      masks base::date()
```

```
## x dplyr::filter()        masks stats::filter()
```

```
## x tsibble::intersect()   masks base::intersect()
```

```
## x lubridate::interval() masks tsibble::interval()
```

```
## x dplyr::lag()           masks stats::lag()
```

```
## x tsibble::setdiff()     masks base::setdiff()
```

```
## x tsibble::union()       masks base::union()
```

```
library(seasonal)
```

```
## Warning: package 'seasonal' was built under R version 4.1.3
```

```
##
```

```
## Attaching package: 'seasonal'
```

```
## The following object is masked from 'package:tibble':
```

```
##
```

```
##      view
```

```
library(glue)
```

```
library(GGally)
```

```
## Warning: package 'GGally' was built under R version 4.1.3
```

```
## Registered S3 method overwritten by 'GGally':
```

```
##   method from
```

```
##   +.gg      ggplot2
```

## 2.10 exercise 9

```
# "Total Private" Employed from us_employment
```

```
employed_total_private <- us_employment %>%
```

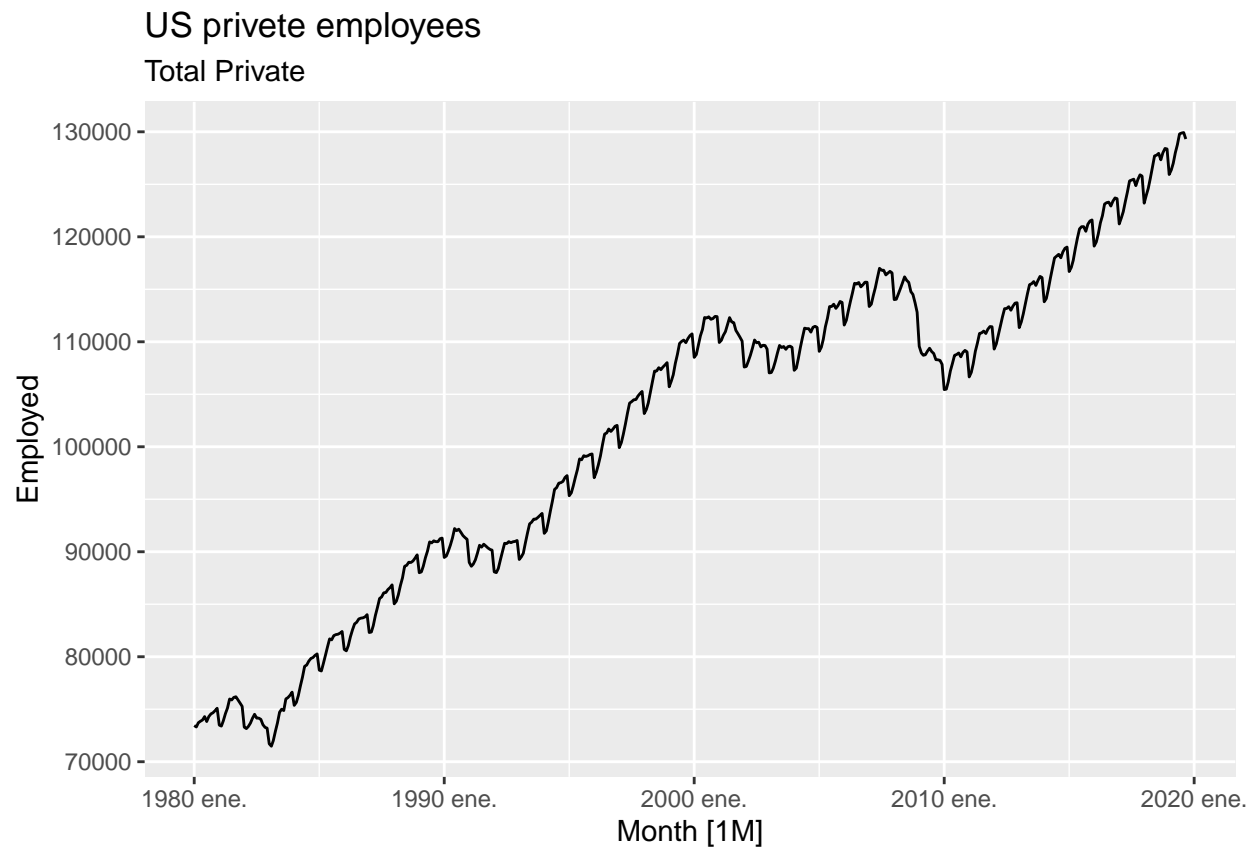
```
  filter(Title == "Total Private" & year(Month) >= 1980) %>%
```

```
  select(Month, Employed)
```

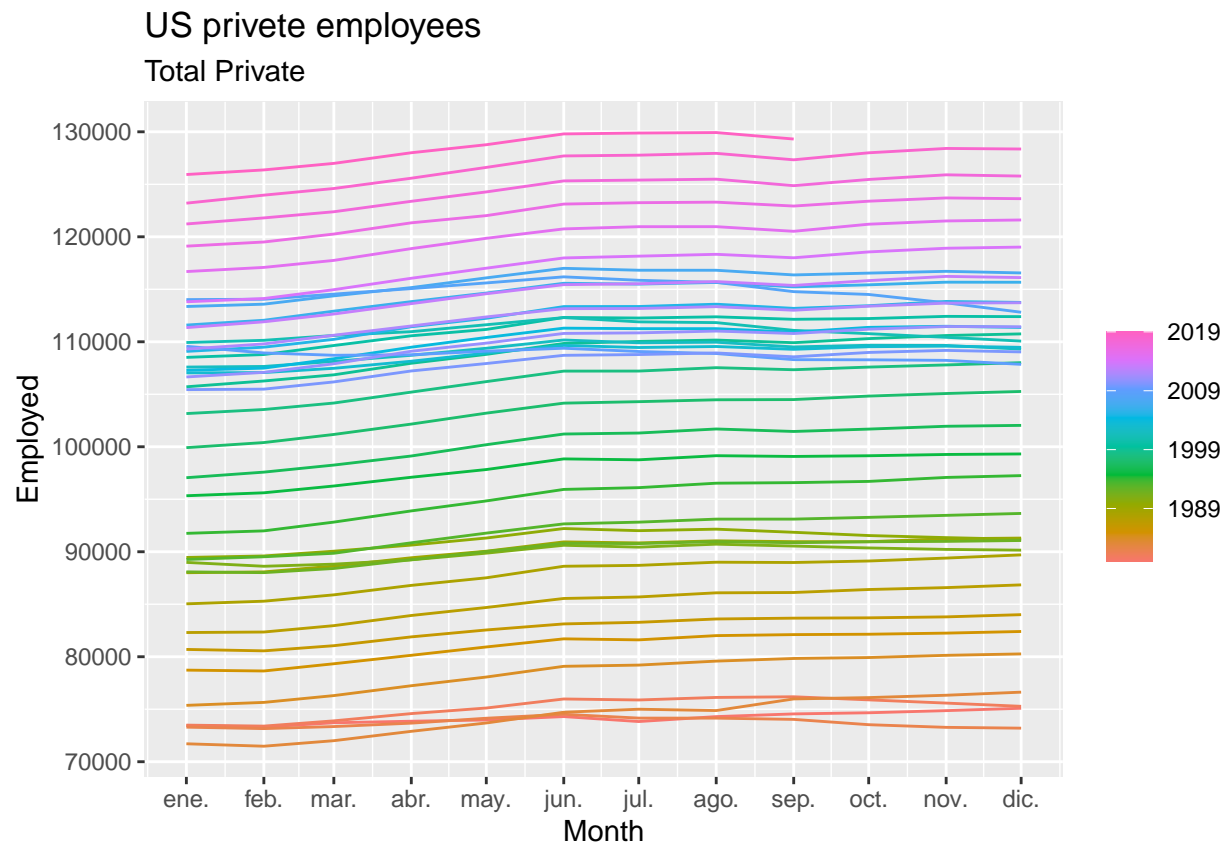
```
autoplot(employed_total_private, Employed) +
```

```
  labs(title = "US private employees",
```

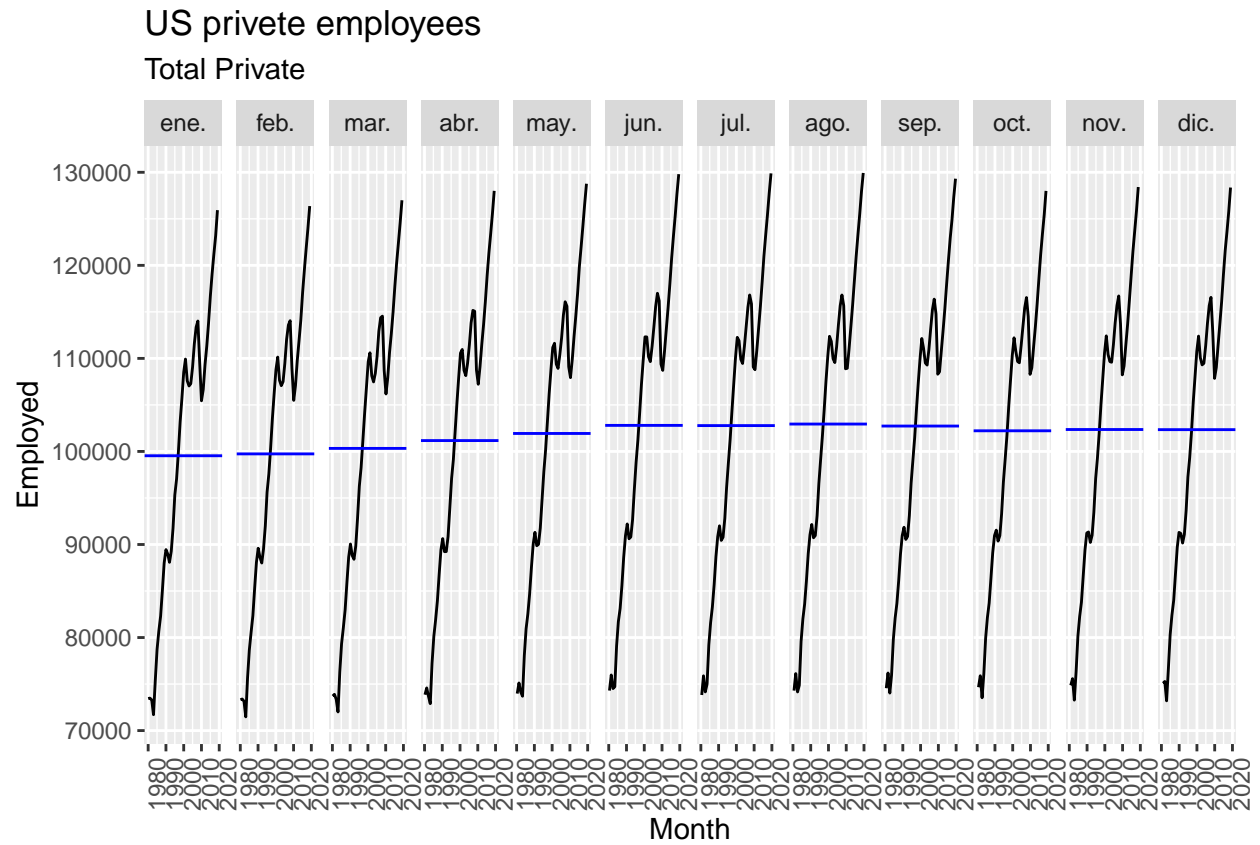
```
        subtitle = "Total Private")
```



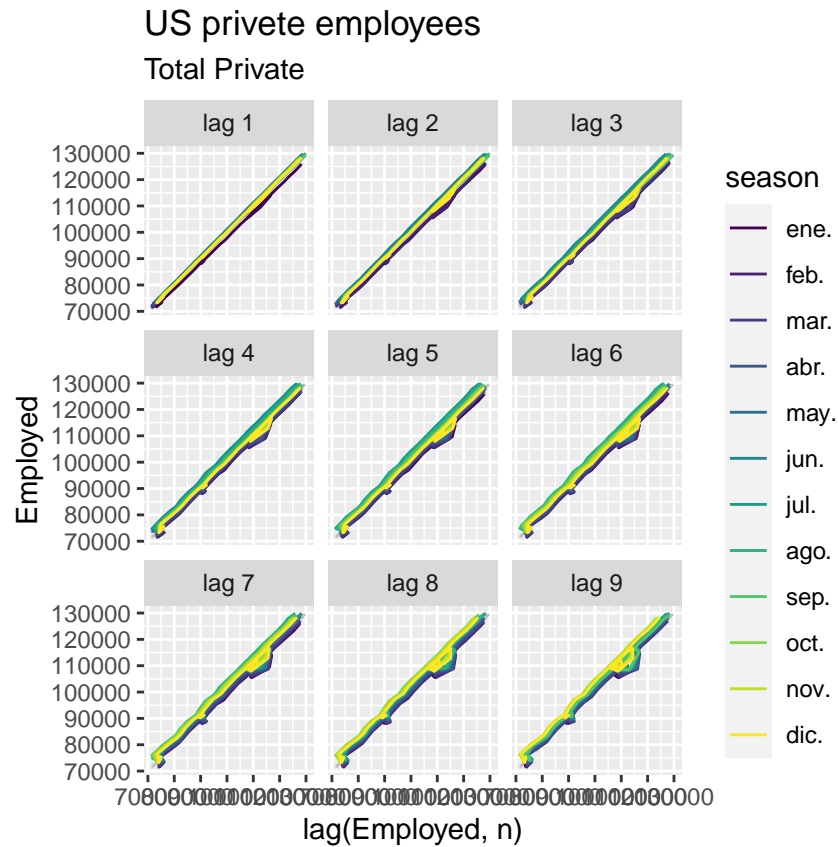
```
gg_season(employed_total_private, Employed) +  
  labs(title = "US private employees",  
        subtitle = "Total Private")
```



```
gg_subseries(employed_total_private, Employed) +
  labs(title = "US private employees",
        subtitle = "Total Private")
```



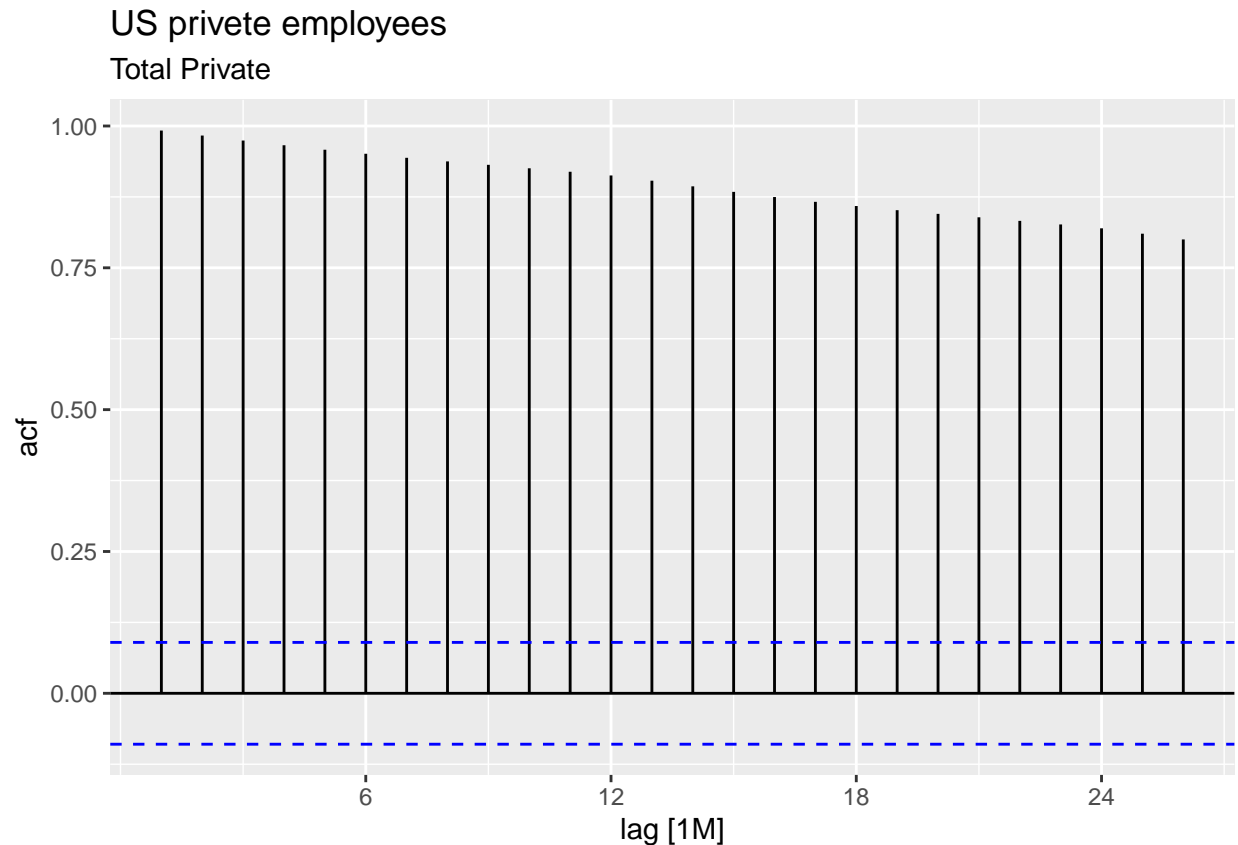
```
gg_lag(employed_total_private, Employed) +
  labs(title = "US private employees",
        subtitle = "Total Private")
```



```
employed_total_private %>% ACF(Employed, lag_max = 9)
```

```
## # A tibble: 9 x 2 [1M]
##   lag   acf
##   <lag> <dbl>
## 1 1M 0.992
## 2 2M 0.983
## 3 3M 0.974
## 4 4M 0.966
## 5 5M 0.958
## 6 6M 0.951
## 7 7M 0.944
## 8 8M 0.938
## 9 9M 0.932
```

```
employed_total_private %>%
  ACF(Employed) %>%
  autoplot() +
  labs(title="US private employees",
       subtitle = "Total Private")
```



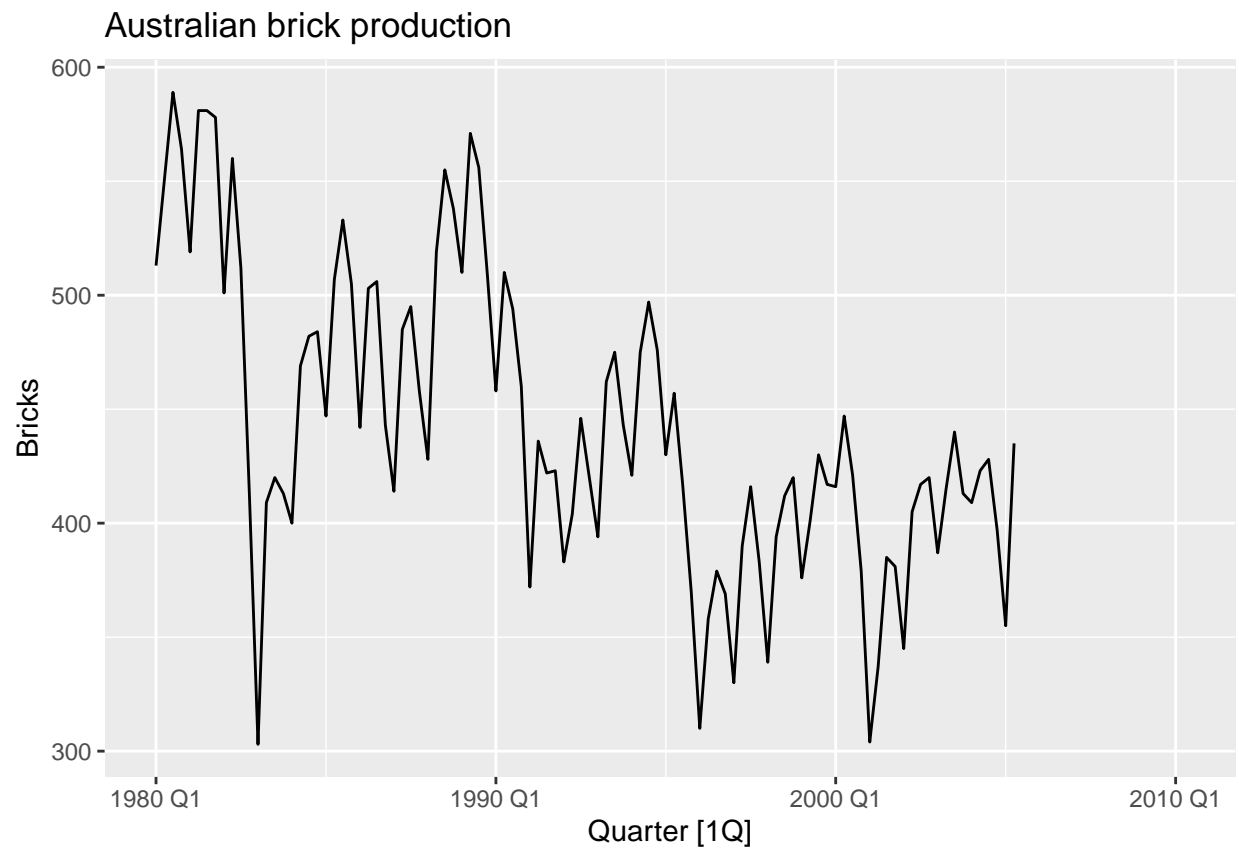
We can see that there is seasonality on the data, the number of employees start as a lower value in January and slowly grow over the month until it gets a peak in the summer month, and then it starts decreasing again. We can see that there is no cyclicity on the data, as there are no longer cycles. The trend is increasing over time, except for the 2007-2008 crises where the employment decrease.

We see an increasing series that grow over time as the country grows, there is an important seasonal pattern over the months, where the employment grows during the first half of the year and decreases over the second half. There are a couple of unusual years during the early 80 decade, the early 90 decade, the early and late 2020 decade, where employment decrease, as it was periods of economic crisis.

```
# Bricks from aus_production
bricks_aus_production <- aus_production %>%
  filter(year(Quarter) >= 1980) %>% #& Title == "Total Private"
  select(Quarter, Bricks)

autoplot(bricks_aus_production, Bricks) +
  labs(title = "Australian brick production")
```

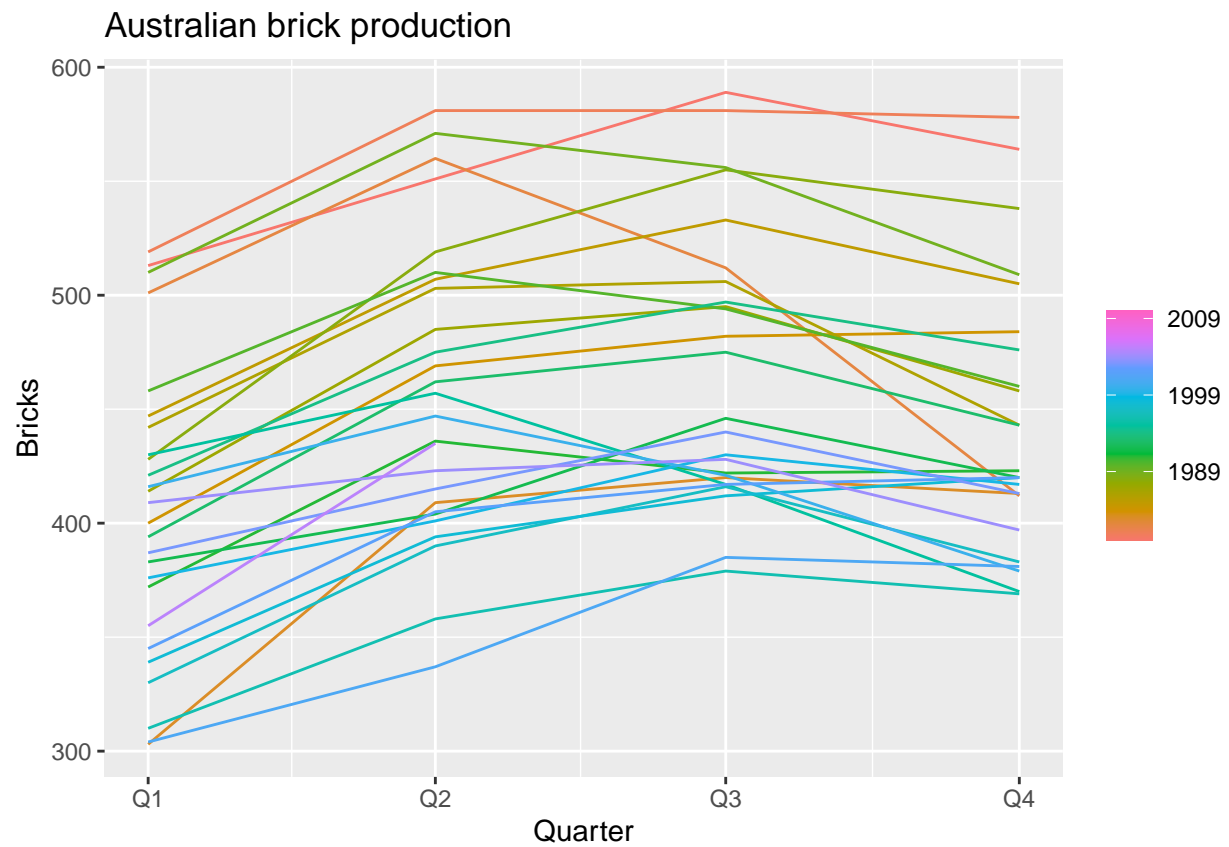
```
## Warning: Removed 20 row(s) containing missing values (geom_path).
```



```
gg_season(bricks_australia_production, Bricks) +  
  labs(title = "Australian brick production")
```


```
## Warning: Removed 20 row(s) containing missing values (geom_path).
```





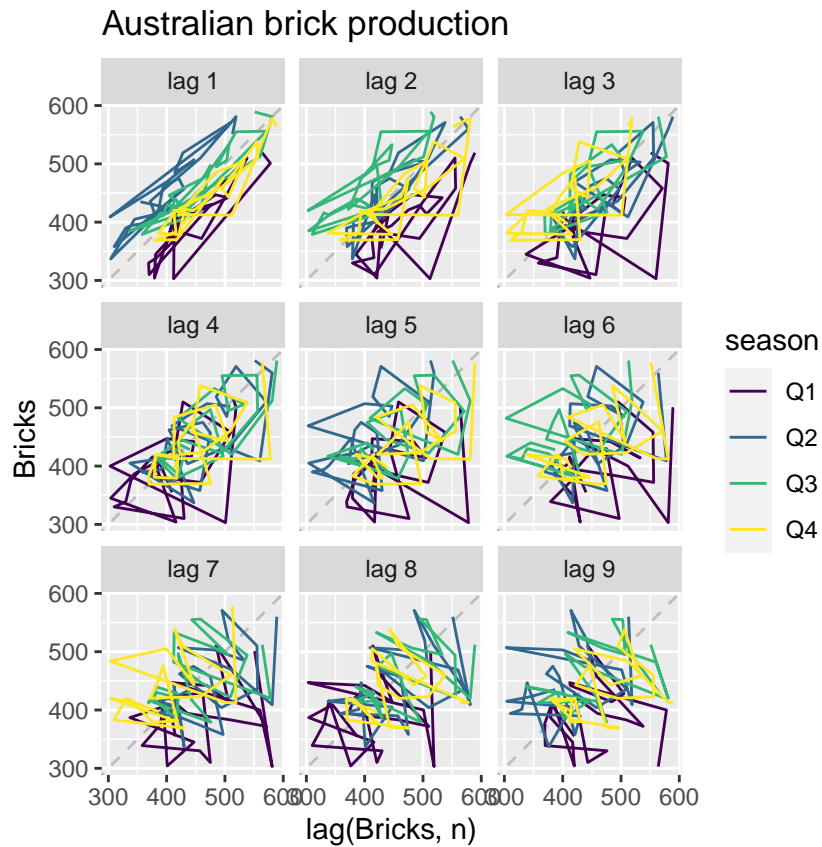
```
gg_subseries(bricks_australia_production, Bricks) +
  labs(title = "Australian brick production")
```

```
## Warning: Removed 5 row(s) containing missing values (geom_path).
```



```
gg_lag(bricks_australia_production, Bricks) +  
  labs(title = "Australian brick production")
```

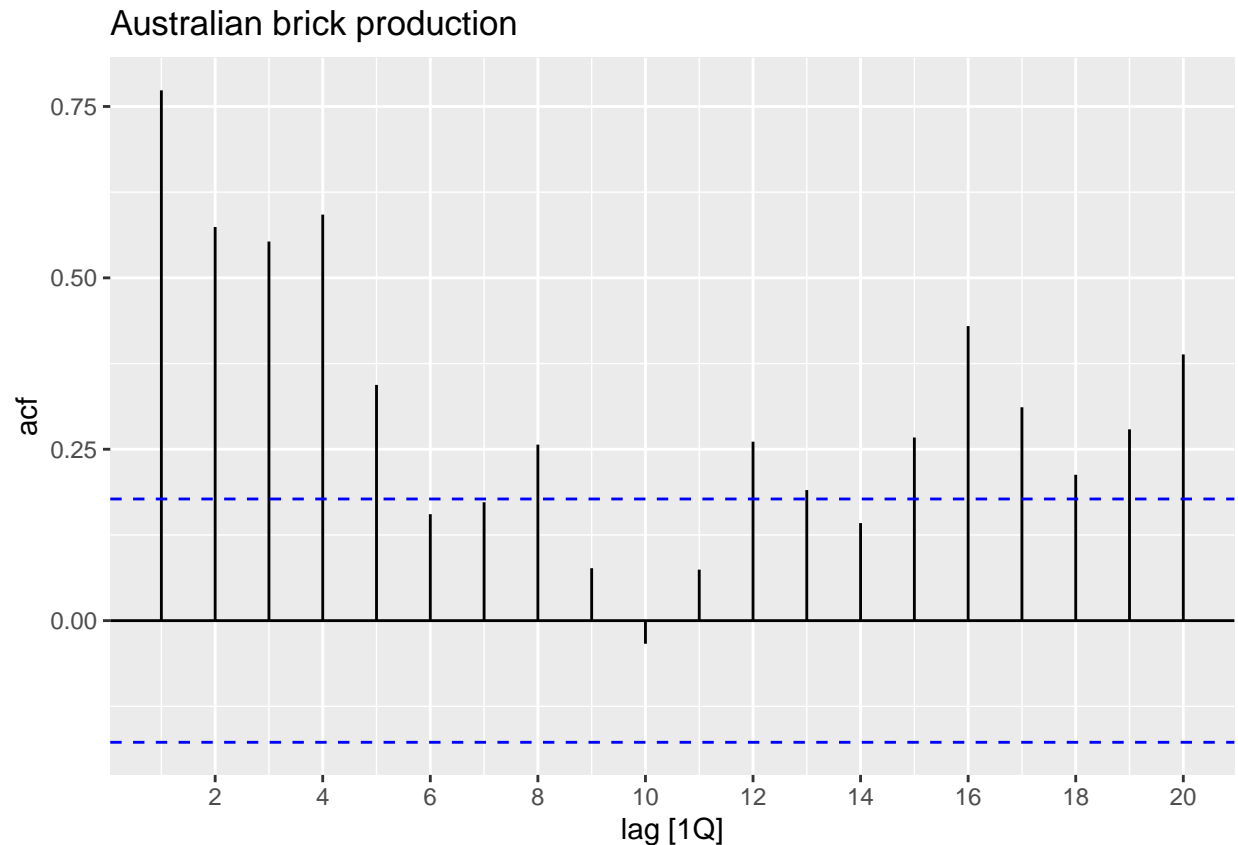
```
## Warning: Removed 20 rows containing missing values (gg_lag).
```



```
bricks_aus_production %>% ACF(Bricks, lag_max = 9)
```

```
## # A tibble: 9 x 2 [1Q]
##   lag   acf
##   <lag> <dbl>
## 1    1Q 0.774
## 2    2Q 0.574
## 3    3Q 0.553
## 4    4Q 0.592
## 5    5Q 0.344
## 6    6Q 0.155
## 7    7Q 0.173
## 8    8Q 0.257
## 9    9Q 0.0765
```


```
bricks_aus_production %>%
  ACF(Bricks) %>%
  autoplot() +
  labs(title="Australian brick production")
```




We can see that there is also seasonality, as the bricks production increases over the quarters, and start decreasing in the last quarter. We can see that there is a decreasing trend in the data. From the series we can see that the bricks production is greater over the middle of the year with respect to the start and the end. We can see that there is an unusual year that the first half of the 80's, as the production dropped to 300, from the 450-550 value it is supposed to be in.

```
# Hare from pelt
hare_production <- pelt %>%
  #filter(Year >= 1980) %>%
  select(Year, Hare)

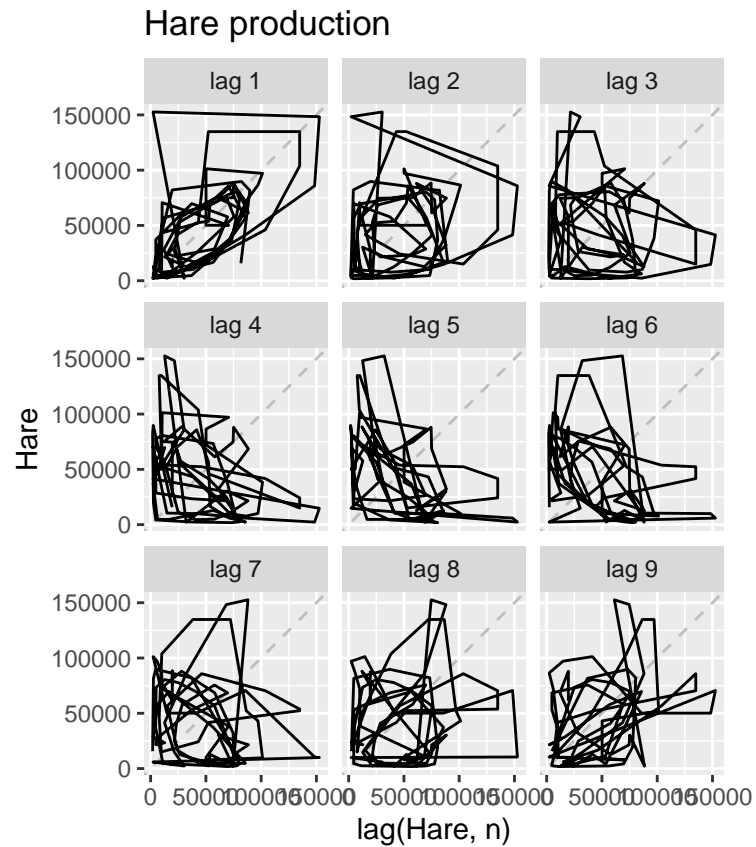
autoplot(hare_production, Hare) +
  labs(title = "Hare production")
```



```
#gg_season(hare_production, Hare) +  
# labs(title = "Hare production")  
  
gg_subseries(hare_production, Hare) +  
  labs(title = "Hare production")
```



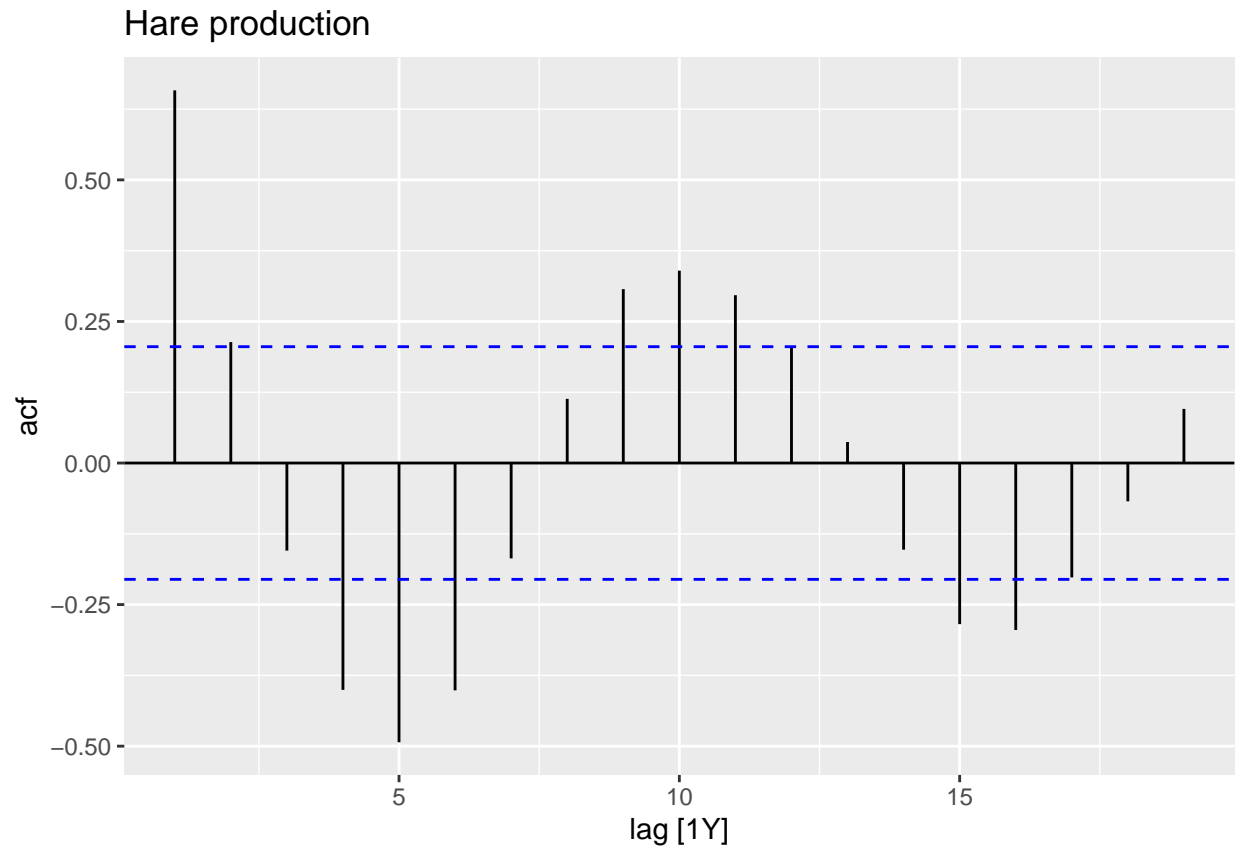
```
gg_lag(hare_production, Hare) +  
  labs(title = "Hare production")
```



```
hare_production %>% ACF(Hare, lag_max = 9)
```

```
## # A tibble: 9 x 2 [1Y]
##   lag   acf
##   <lag> <dbl>
## 1  1Y  0.658
## 2  2Y  0.214
## 3  3Y -0.155
## 4  4Y -0.401
## 5  5Y -0.493
## 6  6Y -0.401
## 7  7Y -0.168
## 8  8Y  0.113
## 9  9Y  0.307
```

```
hare_production %>%
  ACF(Hare) %>%
  autoplot() +
  labs(title="Hare production")
```

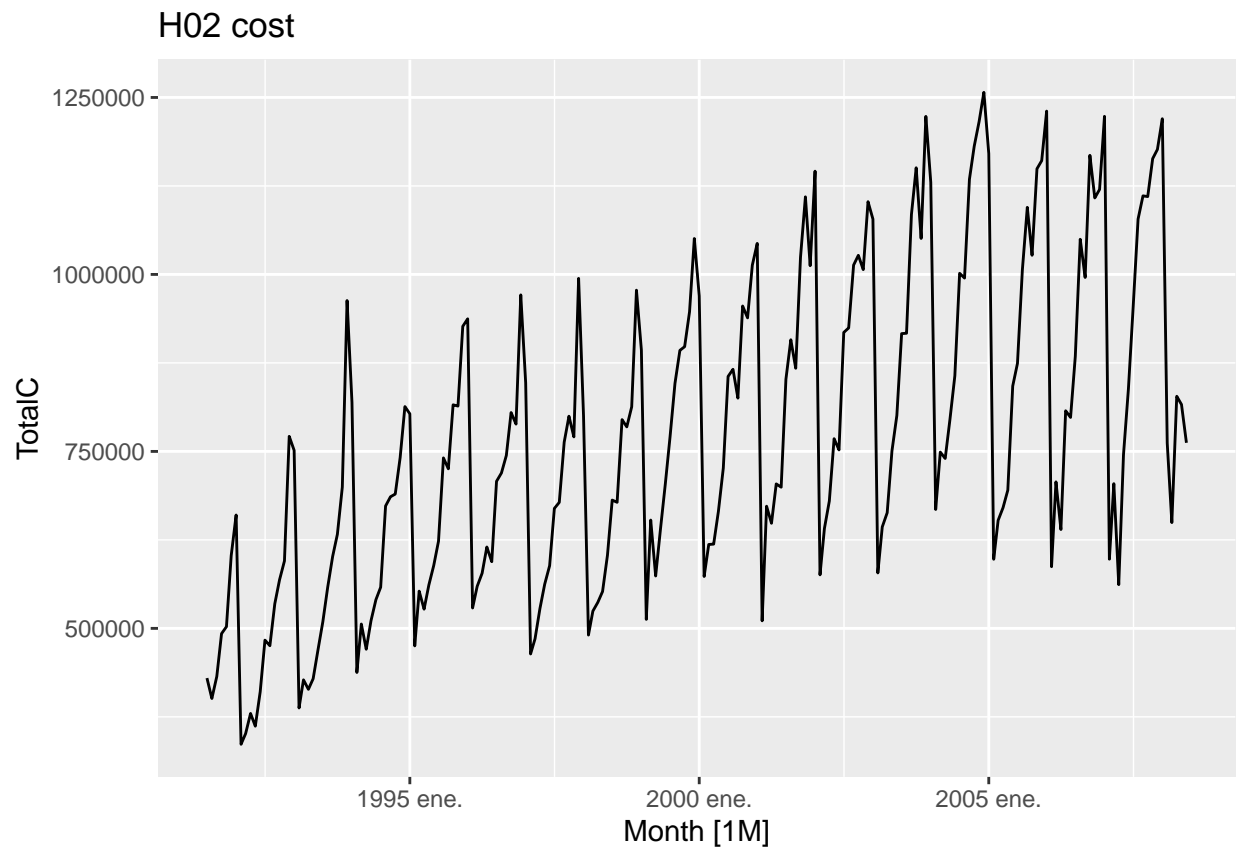


In this series, we only have access to yearly data, so we cannot test if there is any seasonality. Although we can see that there is cyclicity on the data, as the values grows and falls in a period of 5 years, there is also no trend during this period. The only remarkable year is in the early 60's where the data experience the highest value in all the time series.


```
# "H02" Cost from PBS
h02_cost <- PBS %>%
  filter(ATC2 == "H02") %>% #Year >= 1980
  select(Month, Cost) %>%
  summarise(TotalC = sum(Cost))

autoplot(h02_cost, TotalC) +
  labs(title = "H02 cost")
```






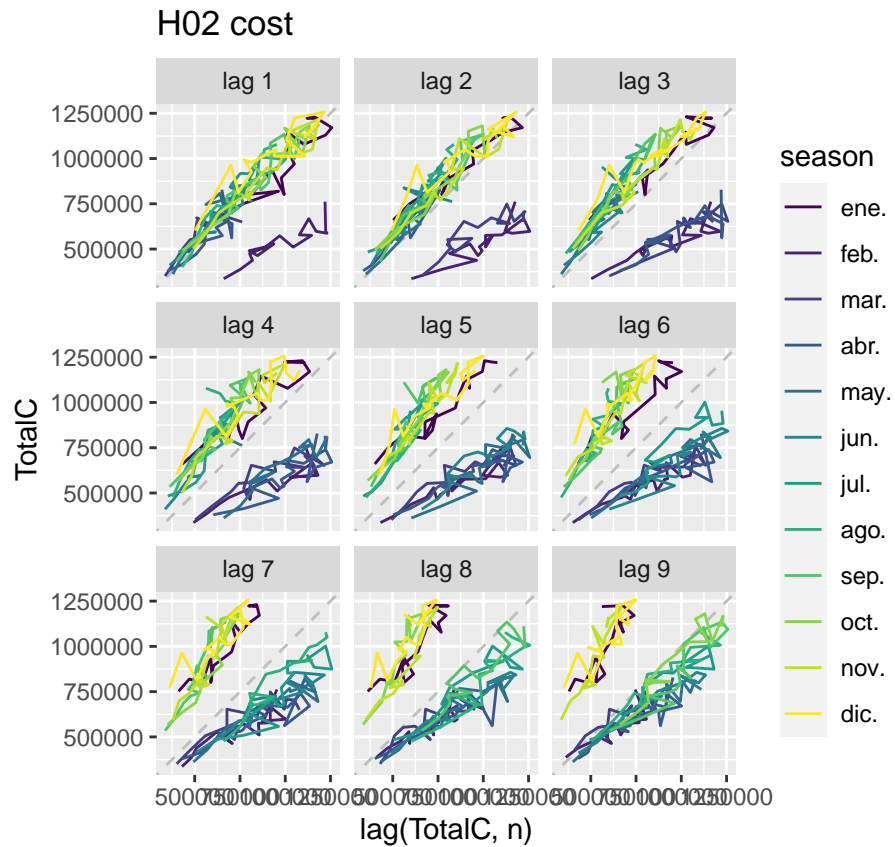
```
gg_season(h02_cost, TotalC) +  
  labs(title = "H02 cost")
```



```
gg_subseries(h02_cost, TotalC) +  
  labs(title = "H02 cost")
```




```
gg_lag(h02_cost, TotalC) +  
  labs(title = "H02 cost")
```



```
h02_cost %>% ACF(TotalC, lag_max = 9)
```

```
## # A tibble: 9 x 2 [1M]
##   lag   acf
##   <lag> <dbl>
## 1 1M 0.755
## 2 2M 0.558
## 3 3M 0.386
## 4 4M 0.228
## 5 5M 0.126
## 6 6M 0.0874
## 7 7M 0.116
## 8 8M 0.203
## 9 9M 0.335
```


```
h02_cost %>%
  ACF(TotalC) %>%
  autoplot() +
  labs(title="H02 cost")
```



We can see that the series has seasonality, it starts as a time high during January, follow by a huge drop in February a slow increase over the rest of the year forward the January values. We can also see that there is no cyclicality, and the general trend is increasing, it is interesting the huge drop on the beginning of the year, this can be related to how medicine is administered in the USA (I'm not so familiarize with how this works). There is no unusual year on the time series.


```
# "H02" Cost from PBS
us_Barrels <- us_gasoline %>%
  #filter(ATC2 == "H02") %>% #Year >= 1980
  #select(Month, Cost) %>%
  #summarise(TotalC = sum(Cost))

autoplot(us_Barrels, Barrels) +
  labs(title = "US Gasoline Barrels")
```

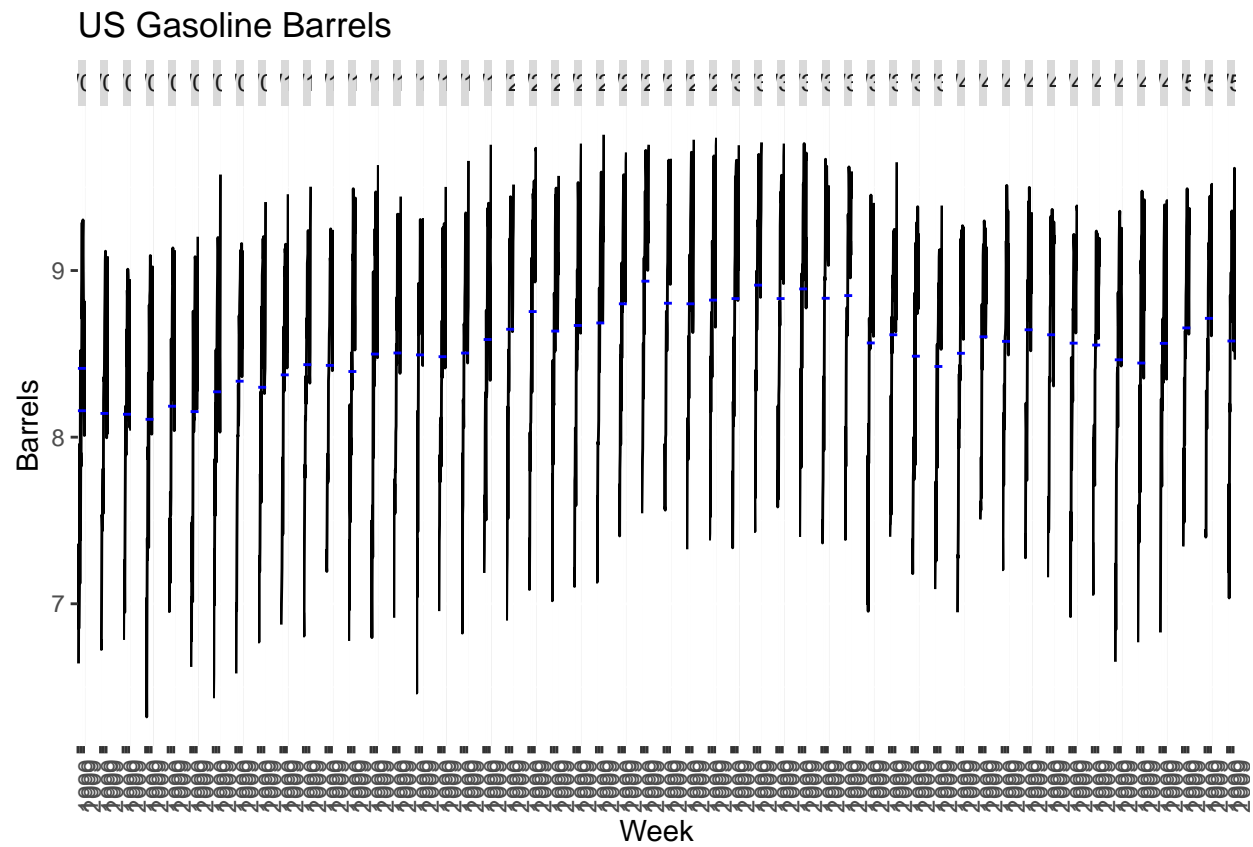


```
gg_season(us_Barrels, Barrels) +  
  labs(title = "US Gasoline Barrels")
```

US Gasoline Barrels

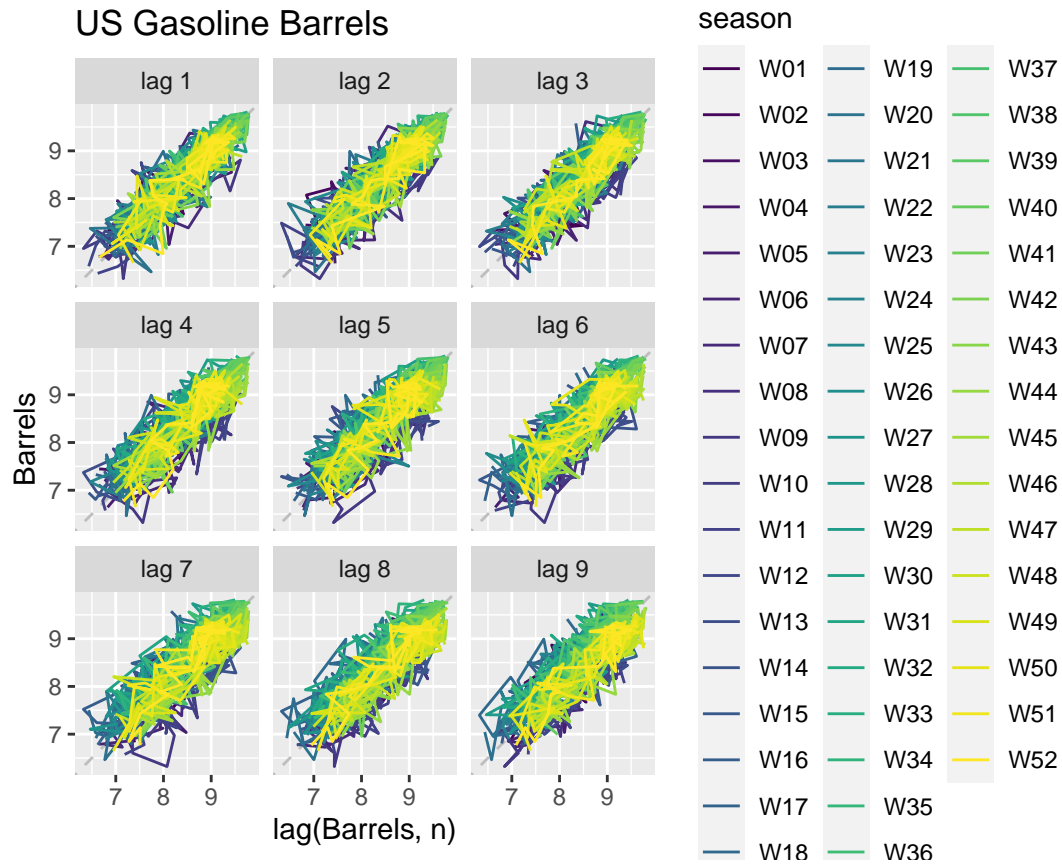


```
gg_subseries(us_Barrels, Barrels) +  
  labs(title = "US Gasoline Barrels")
```



```
gg_lag(us_Barrels, Barrels) +
  labs(title = "US Gasoline Barrels")
```

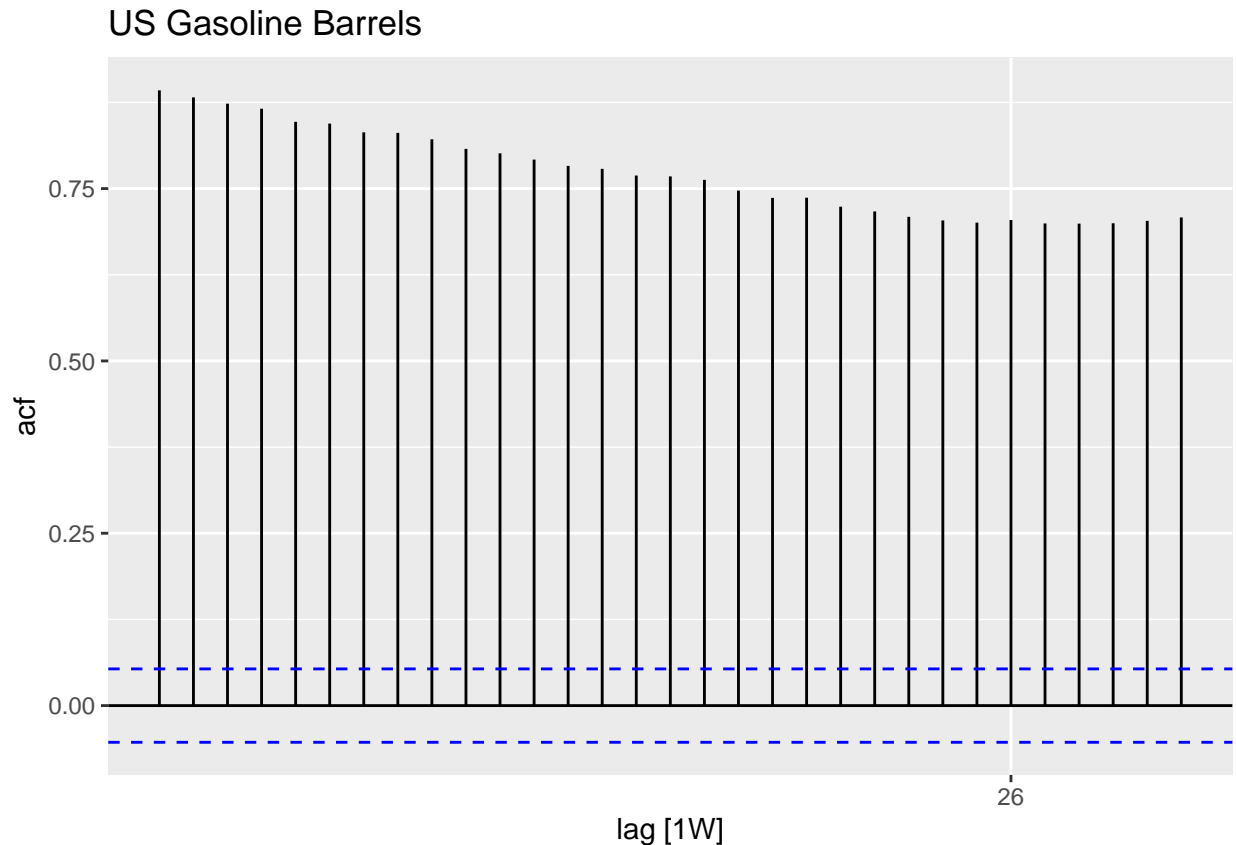




```
us_Barrels %>% ACF(Barrels, lag_max = 9)
```

```
## # A tibble: 9 x 2 [1W]
##   lag   acf
##   <lag> <dbl>
## 1 1W 0.893
## 2 2W 0.882
## 3 3W 0.873
## 4 4W 0.866
## 5 5W 0.847
## 6 6W 0.844
## 7 7W 0.832
## 8 8W 0.831
## 9 9W 0.822
```

```
us_Barrels %>%
  ACF(Barrels) %>%
  autoplot() +
  labs(title="US Gasoline Barrels")
```



In this case there seems to be some seasonality, although it is not extremely clear, we can see, that the barrels increase over the weeks, and then decrease over the end of the year. There is no cyclicality, and there is a positive trend, we can see that the number of barrels increase over time, although in the last part of the series the values stabilize, and the growing pattern is smaller. There is no unusual year on the series.

### 3.7 exercises 9


- a) Here we can see an increasing time series, we can clearly see that the trend is positive in every year. If we take a look at the seasonality of the time series, we can find that there is a seasonal component, in this case it follows a complex path of increasing in values during the first months of the year, followed by a decrease until the month of August, where it reaches a time low, then it increases in September, followed by another fall during October and November, and finally it ends at the highest value during the last month of the year in September.
- b) The recession during 1991/1992 is visible under the remainder component of the decomposition, as we can see a huge drop during that time period, while the normal residual is close to zero, and with a deviation no greater than 100, during that period it drops to -400.

### 3.7 exercises 10


```
# "Total Private" Employed from us_employment
can_gas <- canadian_gas %>%
  #filter(Title == "Total Private" & year(Month) >= 1980) %>%
  #select(Month, Employed)
```

a) Plot the data using `autoplot()`, `gg_subseries()` and `gg_season()` to look at the effect of the changing seasonality over time.


```
autoplot(can_gas, Volume) +  
  labs(title = "Can Gas")
```



```
gg_season(can_gas, Volume) +  
  labs(title = "Can Gas")
```



```
gg_subseries(can_gas, Volume) +  
  labs(title = "Can Gas")
```




We can see there is seasonality, as the values start high at the beginning of the first month of the year, and slowly drop over the month, until the month of June, after that the values start increasing towards the end of the year, this pattern is similar over all the years.

b) Do an STL decomposition of the data. You will need to choose a seasonal window to allow for the changing shape of the seasonal component.

```
can_gas %>%
  model(
    STL(Volume ~ trend(window = 7) +
      season(window = 13),
    robust = TRUE)) %>%
  components() %>%
  autoplot()
```

## STL decomposition


Volume = trend + season\_year + remainder



We can see that the seasonal component starts to increase over the years until it achieves a higher peak in the mid 80's, after then it starts decreasing.

c) How does the seasonal shape change over time? [Hint: Try plotting the seasonal component using `gg_season()`.]

```
gg_season(can_gas, Volume) +  
  labs(title = "Can Gas")
```




We can see that the seasonal component start to get noisier as time advances, at the beginning it was a simple curve, while the last years it had way more jumps and Sharpe edges.

d) Can you produce a plausible seasonally adjusted series?

```
x11_dcmp <- can_gas %>%
  model(x11 = X_13ARIMA_SEATS(Volume ~ x11())) %>%
  components()
autoplot(x11_dcmp) +
  labs(title =
    "Decomposition (x11) of Canadian Gas Volume")
```


## Decomposition (x11) of Canadian Gas Volume

Volume = trend \* seasonal \* irregular



```
x11_dcmp %>%
  ggplot(aes(x = Month)) +
  geom_line(aes(y = Volume, colour = "Data")) +
  geom_line(aes(y = season_adjust,
                colour = "Seasonally Adjusted")) +
  geom_line(aes(y = trend, colour = "Trend")) +
  labs(y = "Volume",
        title = "x11 Canadian Gas Volume") +
  scale_colour_manual(
    values = c("gray", "#0072B2", "#D55E00"),
    breaks = c("Data", "Seasonally Adjusted", "Trend")
  )
```






```
x11_dcmp_seats <- can_gas %>%  
  model(seats = X_13ARIMA_SEATS(Volume ~ seats())) %>%  
  components()  
autoplot(x11_dcmp) +  
  labs(title =  
    "Decomposition (Seats) of Canadian Gas Volume")
```

## Decomposition (Seats) of Canadian Gas Volume

Volume = trend \* seasonal \* irregular



e) Compare the results with those obtained using SEATS and X-11. How are they different? If we take a look at the results, they are pretty much identical, between X-11 and SEATS method


## 4.6 exercise 1

```
pbs_mean <- PBS %>% features(Cost, list(mean = mean)) %>%
  arrange(desc(mean))
#pbs_mean

pbs_HighMean <- PBS %>%
  filter(Concession == pbs_mean$Concession[1] & Type == pbs_mean$Type[1] & ATC1 == pbs_mean$ATC1[1] & A
#pbs_HighMean

autoplot(pbs_HighMean, Cost) +
  labs(title = "PBS Cost of series with highest mean",
        subtitle = paste("(", pbs_mean$Concession[1], ", ", pbs_mean$Type[1], ", ", pbs_mean$ATC1[1], ", ", p
```

## PBS Cost of series with highest mean ( Concessional , Co-payments , C , C10 )




```
pbs_std <- PBS %>% features(Cost, list(sd = sd)) %>%
  filter(sd != 0) %>%
  arrange(sd)
#pbs_std

pbs_LowestStd <- PBS %>%
  filter(Concession == pbs_std$Concession[1] & Type == pbs_std$Type[1] & ATC1 == pbs_std$ATC1[1] & ATC2 == pbs_std$ATC2[1])
#pbs_LowestStd

autoplot(pbs_LowestStd, Cost) +
  labs(title = "PBS Cost of series with Lowest Std",
       subtitle = paste("(", pbs_std$Concession[1], ", ", pbs_std$Type[1], ", ", pbs_std$ATC1[1], ", ", pbs_std$ATC2[1], ")"))
```

## PBS Cost of series with Lowest Std ( General , Co-payments , M , M02 )



## 4.6 exercise 2

```
tourism %>%
  filter(Purpose == "Holiday")
```

```
## # A tsibble: 6,080 x 5 [1Q]
## # Key:   Region, State, Purpose [76]
##   Quarter Region  State      Purpose Trips
##   <qtr> <chr>    <chr>    <chr>    <dbl>
## 1 1998 Q1 Adelaide South Australia Holiday 224.
## 2 1998 Q2 Adelaide South Australia Holiday 130.
## 3 1998 Q3 Adelaide South Australia Holiday 156.
## 4 1998 Q4 Adelaide South Australia Holiday 182.
## 5 1999 Q1 Adelaide South Australia Holiday 185.
## 6 1999 Q2 Adelaide South Australia Holiday 135.
## 7 1999 Q3 Adelaide South Australia Holiday 136.
## 8 1999 Q4 Adelaide South Australia Holiday 169.
## 9 2000 Q1 Adelaide South Australia Holiday 184.
## 10 2000 Q2 Adelaide South Australia Holiday 134.
## # ... with 6,070 more rows
```

```

tourism_features <- tourism %>%
  #filter(Purpose == "Holiday")%>%
  features(Trips, feat_stl)

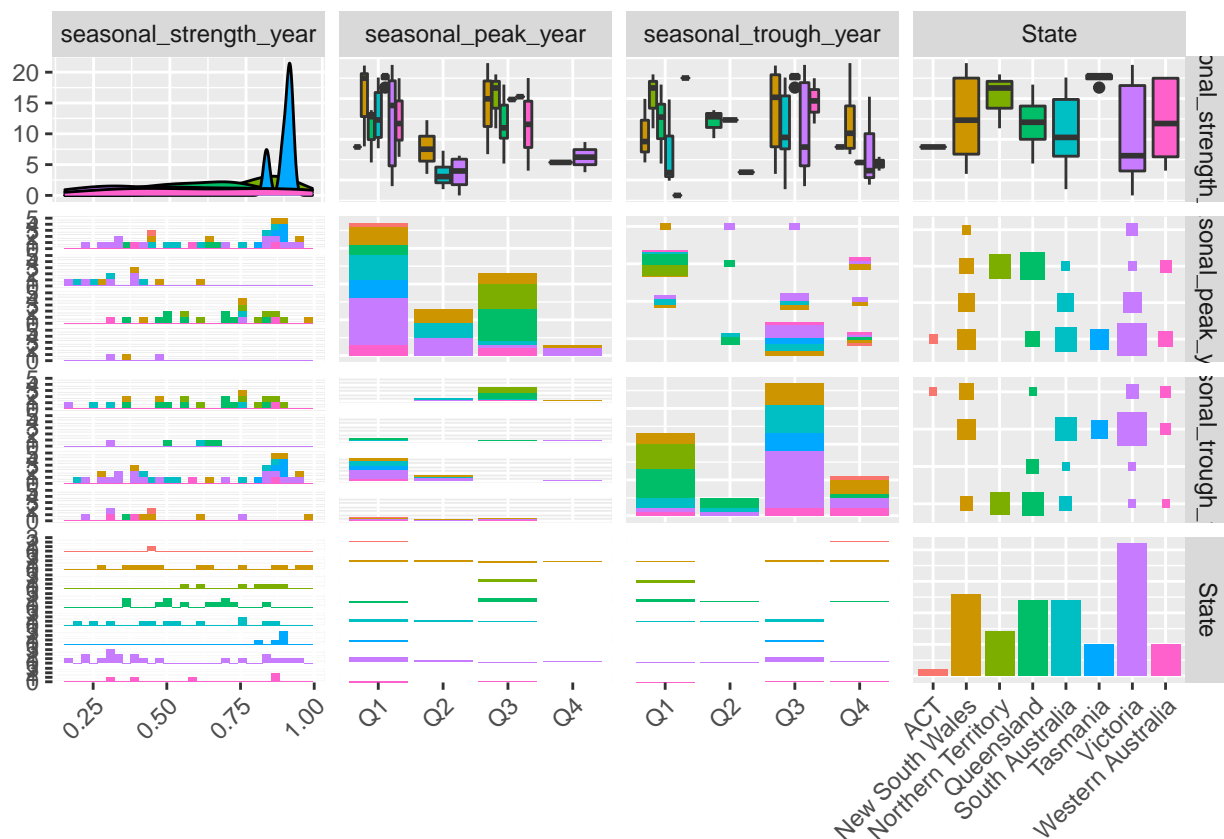
tourism_features %>%
  filter(Purpose == "Holiday") %>%
  select_at(vars(contains("season"), State)) %>%
  mutate(
    seasonal_peak_year = seasonal_peak_year +
      4*(seasonal_peak_year==0),
    seasonal_trough_year = seasonal_trough_year +
      4*(seasonal_trough_year==0),
    seasonal_peak_year = glue("Q{seasonal_peak_year}"),
    seasonal_trough_year = glue("Q{seasonal_trough_year}"),
  ) %>%
  GGally::ggpairs(mapping = aes(colour = State)) +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))

```

```
## Warning: Groups with fewer than two data points have been dropped.
```

```
## Warning in max(ids, na.rm = TRUE): ningun argumento finito para max; retornando
## -Inf
```

```
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
```



We can see that the peak quarter for each state is: ACT: Q1 New South Wales: Q1 Northern Territory: Q3 Queensland: Q3 South Australia: Q1 Tasmania: Q1 Victoria: Q1 Western Australia: Q1

## 4.6 exercise 3

```
PBS_features <- PBS %>%
  #filter(Purpose == "Holiday")%>%
  features(Cost, feat_stl)
PBS_features
```

```
## # A tibble: 336 x 13
##   Concession Type ATC1 ATC2 trend_strength seasonal_streng~ seasonal_peak_y~
##   <chr> <chr> <chr> <chr> <dbl> <dbl> <dbl>
## 1 Concessio~ Co-p~ A A01 0.837 0.899 9
## 2 Concessio~ Co-p~ A A02 0.973 0.937 11
## 3 Concessio~ Co-p~ A A03 0.978 0.845 11
## 4 Concessio~ Co-p~ A A04 0.944 0.843 9
## 5 Concessio~ Co-p~ A A05 0.960 0.870 11
## 6 Concessio~ Co-p~ A A06 0.960 0.951 11
## 7 Concessio~ Co-p~ A A07 0.971 0.910 11
## 8 Concessio~ Co-p~ A A09 0.942 0.891 11
## 9 Concessio~ Co-p~ A A10 0.975 0.917 11
## 10 Concessio~ Co-p~ A A11 0.982 0.898 11
## # ... with 326 more rows, and 6 more variables: seasonal_trough_year <dbl>,
```

```
## #   spikiness <dbl>, linearity <dbl>, curvature <dbl>, stl_e_acf1 <dbl>,  
## #   stl_e_acf10 <dbl>
```

```
PBS_features %>%  
  select_at(vars(contains("season"))) %>%  
  mutate(  
    seasonal_peak_year = seasonal_peak_year +  
      4*(seasonal_peak_year==0),  
    seasonal_trough_year = seasonal_trough_year +  
      4*(seasonal_trough_year==0),  
    seasonal_peak_year = glue("Q{seasonal_peak_year}"),  
    seasonal_trough_year = glue("Q{seasonal_trough_year}"),  
  ) %>%  
  GGally::ggpairs(mapping = aes()) +  
  theme(axis.text.x = element_text(angle = 45, hjust = 1))
```

```
## Warning: Removed 2 rows containing non-finite values (stat_density).
```

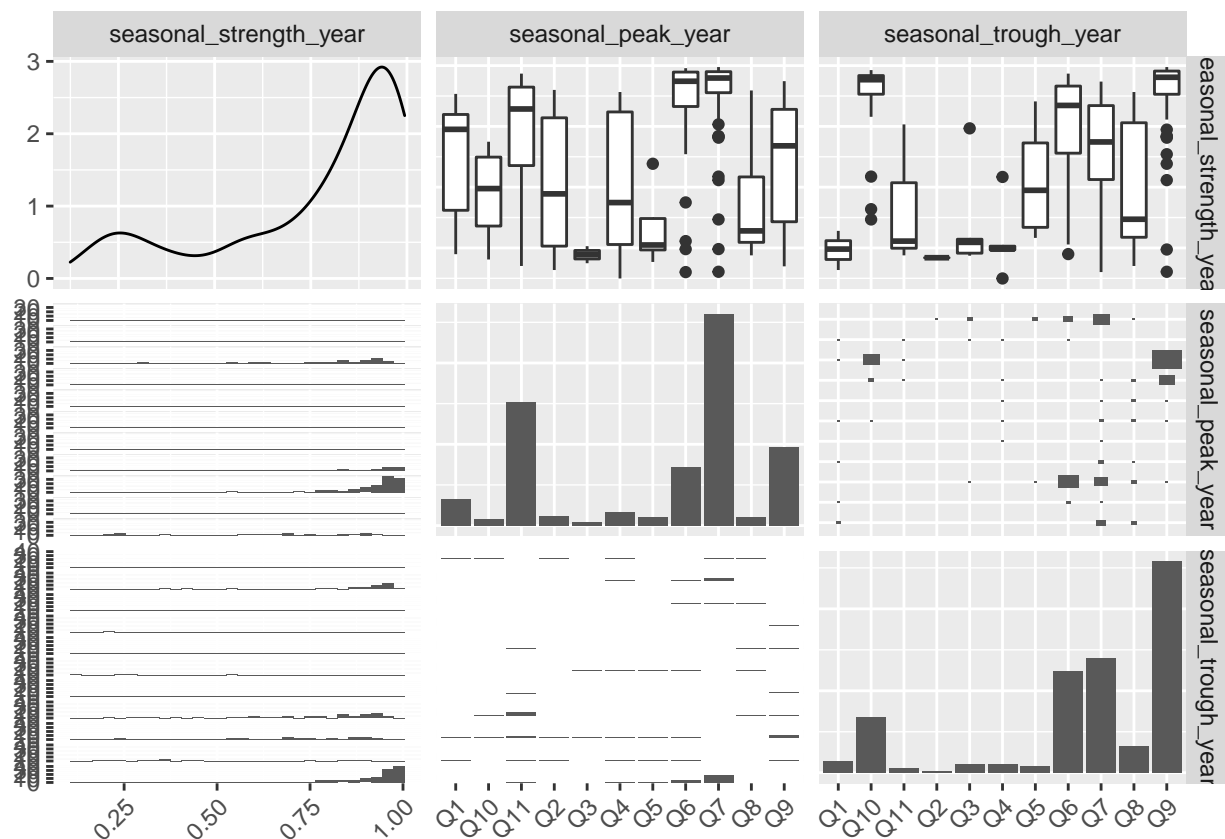
```
## Warning: Removed 2 rows containing non-finite values (stat_boxplot).  
## Removed 2 rows containing non-finite values (stat_boxplot).
```

```
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
```

```
## Warning: Removed 2 rows containing non-finite values (stat_bin).
```

```
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
```

```
## Warning: Removed 2 rows containing non-finite values (stat_bin).
```



## 5.11 exercises 9

a) Create a training set for household wealth (hh\_budget) by withholding the last four years as a test set.

hh\_budget

```
## # A tibble: 88 x 8 [1Y]
## # Key:   Country [4]
##   Country Year Debt   DI Expenditure Savings Wealth Unemployment
##   <chr>   <dbl> <dbl> <dbl>      <dbl>      <dbl> <dbl>      <dbl>
## 1 Australia 1995  95.7 3.72      3.40      5.24   315.      8.47
## 2 Australia 1996  99.5 3.98      2.97      6.47   315.      8.51
## 3 Australia 1997 108. 2.52      4.95      3.74   323.      8.36
## 4 Australia 1998 115. 4.02      5.73      1.29   339.      7.68
## 5 Australia 1999 121. 3.84      4.26      0.638  354.      6.87
## 6 Australia 2000 126. 3.77      3.18      1.99   350.      6.29
## 7 Australia 2001 132. 4.36      3.10      3.24   348.      6.74
## 8 Australia 2002 149. 0.0218    4.03     -1.15   349.      6.37
## 9 Australia 2003 159. 6.06      5.04     -0.413  360.      5.93
## 10 Australia 2004 170. 5.53      4.54      0.657  379.      5.40
## # ... with 78 more rows
```



```
train <- hh_budget %>%
  filter(Year <= 2012 & Country=="USA") #<del> Country=="_" | Australia, Canada, Japan, USA
train
```


```
## # A tsibble: 18 x 8 [1Y]
## # Key:      Country [1]
##   Country Year Debt      DI Expenditure Savings Wealth Unemployment
##   <chr>   <dbl> <dbl>   <dbl>      <dbl>   <dbl>   <dbl>      <dbl>
## 1 USA     1995  94.3  3.00      2.95     7.24    472.      5.61
## 2 USA     1996  96.2  2.86      3.47     6.79    485.      5.42
## 3 USA     1997  97.3  3.46      3.77     6.58    510.      4.95
## 4 USA     1998  98.5  5.66      5.31     7.06    527.      4.51
## 5 USA     1999 103.   3.20      5.27     5.27    563.      4.22
## 6 USA     2000 104.   4.69      5.08     5.03    521.      3.99
## 7 USA     2001 108.   2.82      2.52     5.25    499.      4.73
## 8 USA     2002 113.   3.29      2.57     6.06    461.      5.78
## 9 USA     2003 121.   2.79      3.18     5.75    492.      5.99
## 10 USA    2004 128.   3.28      3.75     5.35    525.      5.53
## 11 USA    2005 136.   1.34      3.56     3.28    550.      5.07
## 12 USA    2006 141.   3.56      3.06     4.00    563.      4.62
## 13 USA    2007 144.   1.78      2.22     3.88    561.      4.62
## 14 USA    2008 137.   1.74     -0.212    5.18    479.      5.78
## 15 USA    2009 136.  -0.0516  -1.25     6.34    500.      9.27
## 16 USA    2010 128.   1.01      1.75     6.78    520.      9.62
## 17 USA    2011 121.   2.23      1.89     7.40    503.      8.95
## 18 USA    2012 115.   2.96      1.50     9.14    514.      8.07
```

b) Fit all the appropriate benchmark methods to the training set and forecast the periods covered by the test set.

```
# Fit the models
budget_fit <- train %>%
  model(
    Mean = MEAN(Wealth),
    `Naïve` = NAIVE(Wealth),
    Drift = RW(Wealth ~ drift())
  )

# Generate forecasts for 14 quarters
budget_fc <- budget_fit %>% forecast(h = 4)

#Plot
budget_fc %>%
  autoplot(hh_budget, level = NULL) +
  labs(y = "$US",
    title = "household Wealth USA",
    subtitle = "(1995 - 2013) + 4 year forecast") +
  guides(colour = guide_legend(title = "Forecast"))
```



c) Compute the accuracy of your forecasts. Which method does best?

```
accuracy(budget_fc, hh_budget)
```

```
## # A tibble: 3 x 11
##   .model Country .type    ME  RMSE   MAE   MPE  MAPE  MASE  RMSSE  ACF1
##   <chr>   <chr>   <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 Drift   USA     Test  75.9  76.2  75.9  12.7  12.7  2.88  2.43 -0.561
## 2 Mean   USA     Test  82.9  83.3  82.9  13.9  13.9  3.15  2.65 -0.423
## 3 Naïve  USA     Test  82.1  82.5  82.1  13.8  13.8  3.12  2.63 -0.423
```

Drift method achieve better metrics than the other methods


d) Do the residuals from the best method resemble white noise?

```
train %>%
  model(Drift = RW(Wealth ~ drift())) %>%
  gg_tsresiduals() +
  labs(title = "Residuals from the Drift method")
```

```
## Warning: Removed 1 row(s) containing missing values (geom_path).
```

```
## Warning: Removed 1 rows containing missing values (geom_point).
```

```
## Warning: Removed 1 rows containing non-finite values (stat_bin).
```



Yes, they resemble white noise, as they are concentrated close to the 0 value, although they seem to be a little skewed to the positive size. From the ACF of the residuals we can see that there is no correlation, meaning that the forecast is good.

## 5.11 exercises 10

a) Create a training set for Australian takeaway food turnover (`aus_retail`) by withholding the last four years as a test set.

```
#aus_retail
test_set <- aus_retail %>%
  filter(Industry=="Takeaway food services") %>%
  summarise(AVGTurnover = mean(Turnover))
test_set
```

```
## # A tibble: 441 x 2 [1M]
##   Month AVGTurnover
##   <mtm>      <dbl>
## 1 1982 abr.      27.7
## 2 1982 may.      27.7
## 3 1982 jun.      26.6
## 4 1982 jul.      27.1
## 5 1982 ago.      27.2
## 6 1982 sep.      27.9
## 7 1982 oct.      29.9
```

```
## 8 1982 nov.      30.4
## 9 1982 dic.      34.0
## 10 1983 ene.     32.1
## # ... with 431 more rows
```

```
train <- test_set %>%
  filter(year(Month) <= 2014)

train
```

```
## # A tibble: 393 x 2 [1M]
##       Month AVGTurnover
##       <mtm>      <dbl>
## 1 1982 abr.      27.7
## 2 1982 may.      27.7
## 3 1982 jun.      26.6
## 4 1982 jul.      27.1
## 5 1982 ago.      27.2
## 6 1982 sep.      27.9
## 7 1982 oct.      29.9
## 8 1982 nov.      30.4
## 9 1982 dic.      34.0
## 10 1983 ene.     32.1
## # ... with 383 more rows
```

b) Fit all the appropriate benchmark methods to the training set and forecast the periods covered by the test set.


```
# Fit the models
retail_fit <- train %>%
  model(
    Mean = MEAN(AVGTurnover),
    `Naïve` = NAIVE(AVGTurnover),
    Drift = RW(AVGTurnover ~ drift())
  )

# Generate forecasts for 14 quarters
retail_fc <- retail_fit %>% forecast(h = 48) # 4*12

#Plot
retail_fc %>%
  autoplot(test_set, level = NULL) +
  labs(y = "$US",
    title = "Australian takeaway food turnover",
    subtitle = "(1982 abr. - 2014 dic. ) + 4 year forecast") +
  guides(colour = guide_legend(title = "Forecast"))
```

```
## Warning in grid.Call(C_textBounds, as.graphicsAnnot(x$label), x$x, x$y, : font
## width unknown for character 0x9
```

```
## Warning in grid.Call(C_textBounds, as.graphicsAnnot(x$label), x$x, x$y, : font
## width unknown for character 0x9
```

[illegible]

c) Compute the accuracy of your forecasts. Which method does best?

```
accuracy(retail_fc, test_set)
```

```
## # A tibble: 3 x 10
##   .model .type      ME  RMSE   MAE   MPE  MAPE  MASE  RMSSE  ACF1
##   <chr>  <chr>  <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 Drift  Test  -11.5  16.2  13.4  -6.71  7.58  2.54  2.44  0.408
```

```
## 2 Mean Test 103. 104. 103. 55.2 55.2 19.6 15.7 0.613
## 3 Naïve Test -1.55 14.8 12.0 -1.49 6.66 2.29 2.24 0.613
```

Naïve method achieve better metrics than the other methods

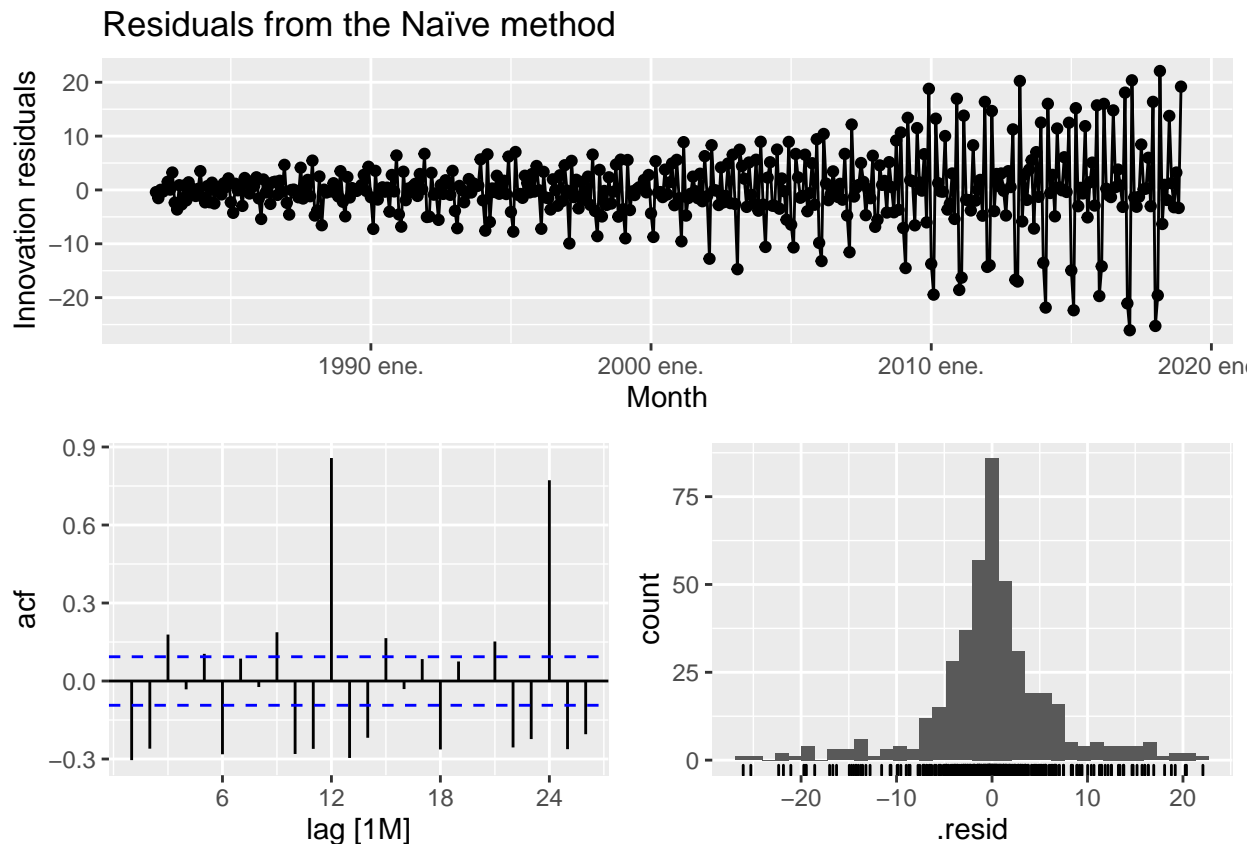
d) Do the residuals from the best method resemble white noise?

```
test_set %>%
  model(Drift = RW(AVGTurnover ~ drift())) %>%
  gg_tsresiduals() +
  labs(title = "Residuals from the Naïve method")
```

```
## Warning: Removed 1 row(s) containing missing values (geom_path).
```

```
## Warning: Removed 1 rows containing missing values (geom_point).
```

```
## Warning: Removed 1 rows containing non-finite values (stat_bin).
```



In this case we can see that the residuals increase over time, although they are grouped by the 0 value, and they resemble a normal distribution. From the ACF of the residuals we can see that there is correlation in multiple values, which means that the model hasn't capture all the information, this is because the naïve model doesn't capture the seasonality of the data.

## 5.11 exercises 11

```
bricks_aus_production <- aus_production %>%
  select(Quarter, Bricks) %>%
  drop_na()
bricks_aus_production
```

```
## # A tsibble: 198 x 2 [1Q]
##   Quarter Bricks
##   <qtr>   <dbl>
## 1 1956 Q1    189
## 2 1956 Q2    204
## 3 1956 Q3    208
## 4 1956 Q4    197
## 5 1957 Q1    187
## 6 1957 Q2    214
## 7 1957 Q3    227
## 8 1957 Q4    222
## 9 1958 Q1    199
## 10 1958 Q2    229
## # ... with 188 more rows
```


a) Use an STL decomposition to calculate the trend-cycle and seasonal indices. (Experiment with having fixed or changing seasonality.)

```
dcmp <- bricks_aus_production %>%
  model(
    STL(Bricks ~ season(window = 13) #"periodic"
    )) %>%
  components()

autoplot(dcmp) +
  labs(title = "STL decomposition ",
       subtitle = "(season window: 13)")
```

## STL decomposition

(season window: 13)




```
dcmp <- bricks_aus_production %>%  
  model(  
    STL(Bricks ~ season(window = 25) #"periodic"  
    )) %>%  
  components()  
  
autoplot(dcmp) +  
  labs(title = "STL decomposition ",  
        subtitle = "(season window: 25)")
```



## STL decomposition


(season window: 25)



```
dcmp <- bricks_aus_production %>%  
  model(  
    STL(Bricks ~ season(window = "periodic") #, robust = TRUE  
  )) %>%  
  components()  
  
autoplot(dcmp) +  
  labs(title = "STL decomposition ",  
       subtitle = "(season window: periodic)")
```

## STL decomposition


(season window: periodic)



b) Compute and plot the seasonally adjusted data.

```
dcmp %>%
  as_tsibble() %>%
  autoplot(Bricks, colour="gray") +
  geom_line(aes(y=season_adjust), colour = "#D55E00") +
  labs(title = "Australian Bricks Production")
```

## Australian Bricks Production




c) Use a naïve method to produce forecasts of the seasonally adjusted data.

```
dcmp2 <- dcmp %>%
  select(-.model)

bricks_fit <- dcmp2 %>%
  model(NAIVE(season_adjust))


# Generate forecasts for 32 quarters
bricks_fc <- bricks_fit %>% forecast(h = 32) # 4*12

#Plot
bricks_fc %>%
  autoplot(dcmp2, level = NULL) +
  labs(y = "$US",
       title = "Australian Bricks season_adjust") +
  guides(colour = guide_legend(title = "Forecast"))
```



d) Use `decomposition_model()` to reseasonalise the results, giving forecasts for the original data.

```
fit_dcmp <- bricks_aus_production %>%
  model(stlf = decomposition_model(
    STL(Bricks ~ season(window = "periodic")), #, robust = TRUE
    NAIVE(season_adjust)
  ))
fit_dcmp %>%
  forecast() %>%
  autoplot(bricks_aus_production)+
  labs(title = "Australian Bricks")
```



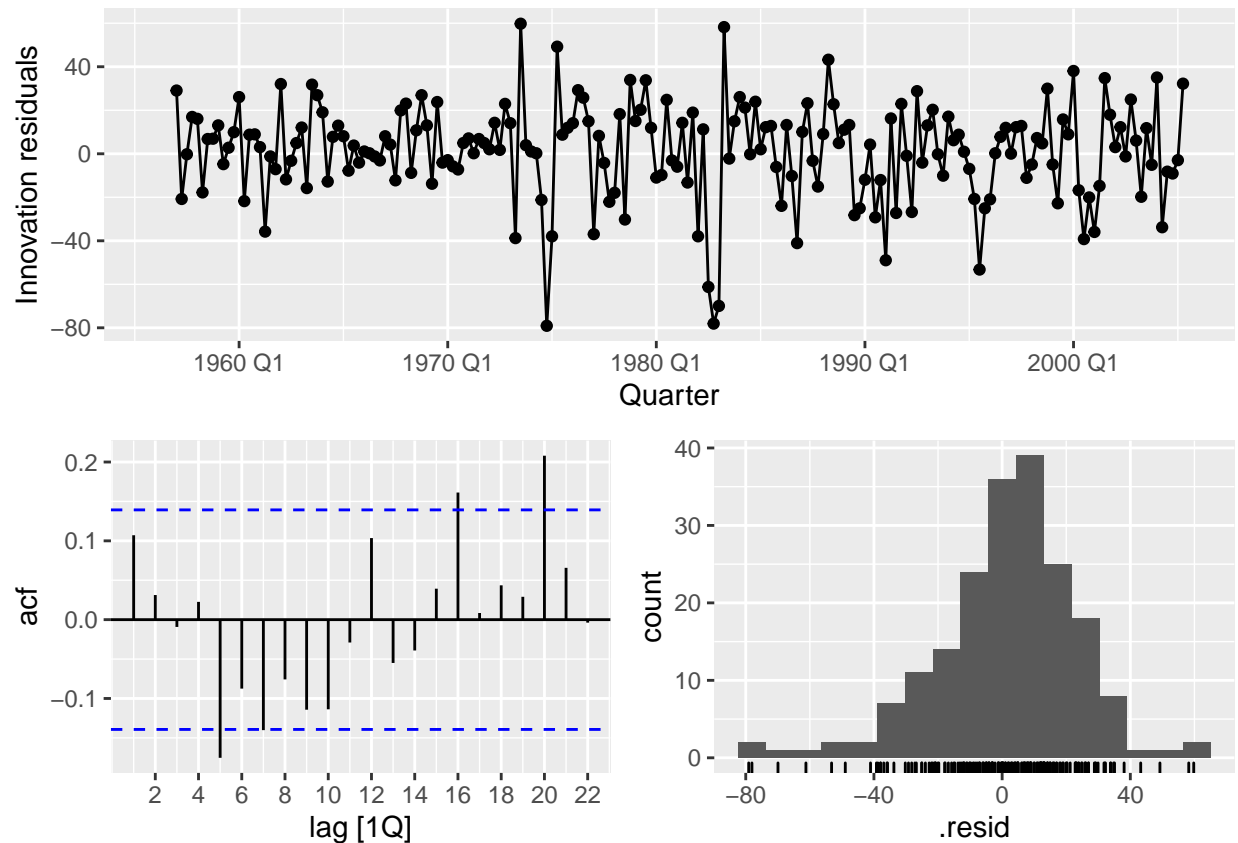
e) Do the residuals look uncorrelated?

```
fit_dcmp %>% gg_tsresiduals()
```

```
## Warning: Removed 4 row(s) containing missing values (geom_path).
```

```
## Warning: Removed 4 rows containing missing values (geom_point).
```

```
## Warning: Removed 4 rows containing non-finite values (stat_bin).
```



There are values lags for which the residuals are correlated, meaning that there is still information remaining on the residuals.


f) Repeat with a robust STL decomposition. Does it make much difference?

```
dcmp <- bricks_aus_production %>%
  model(
    STL(Bricks ~ season(window = "periodic"), #"periodic"
    robust = TRUE)) %>%
  components()

autoplot(dcmp) +
  labs(title = "STL decomposition ",
       subtitle = "(season window: periodic)")
```


## STL decomposition

(season window: periodic)



```
dcmp %>%
  as_tsibble() %>%
  autoplot(Bricks, colour="gray") +
  geom_line(aes(y=season_adjust), colour = "#D55E00") +
  labs(title = "Australian Bricks Production")
```

## Australian Bricks Production




```
dcmp2 <- dcmp %>%
  select(-.model)

bricks_fit <- dcmp2 %>%
  model(NAIVE(season_adjust))

# Generate forecasts for 32 quarters
bricks_fc <- bricks_fit %>% forecast(h = 32) # 4*12

fit_dcmp <- bricks_au_production %>%
  model(stlf = decomposition_model(
    STL(Bricks ~ season(window = "periodic"), robust = TRUE), #, robust = TRUE
    NAIVE(season_adjust)
  ))
fit_dcmp %>%
  forecast() %>%
  autoplot(bricks_au_production)+
  labs(title = "Australian Bricks")
```



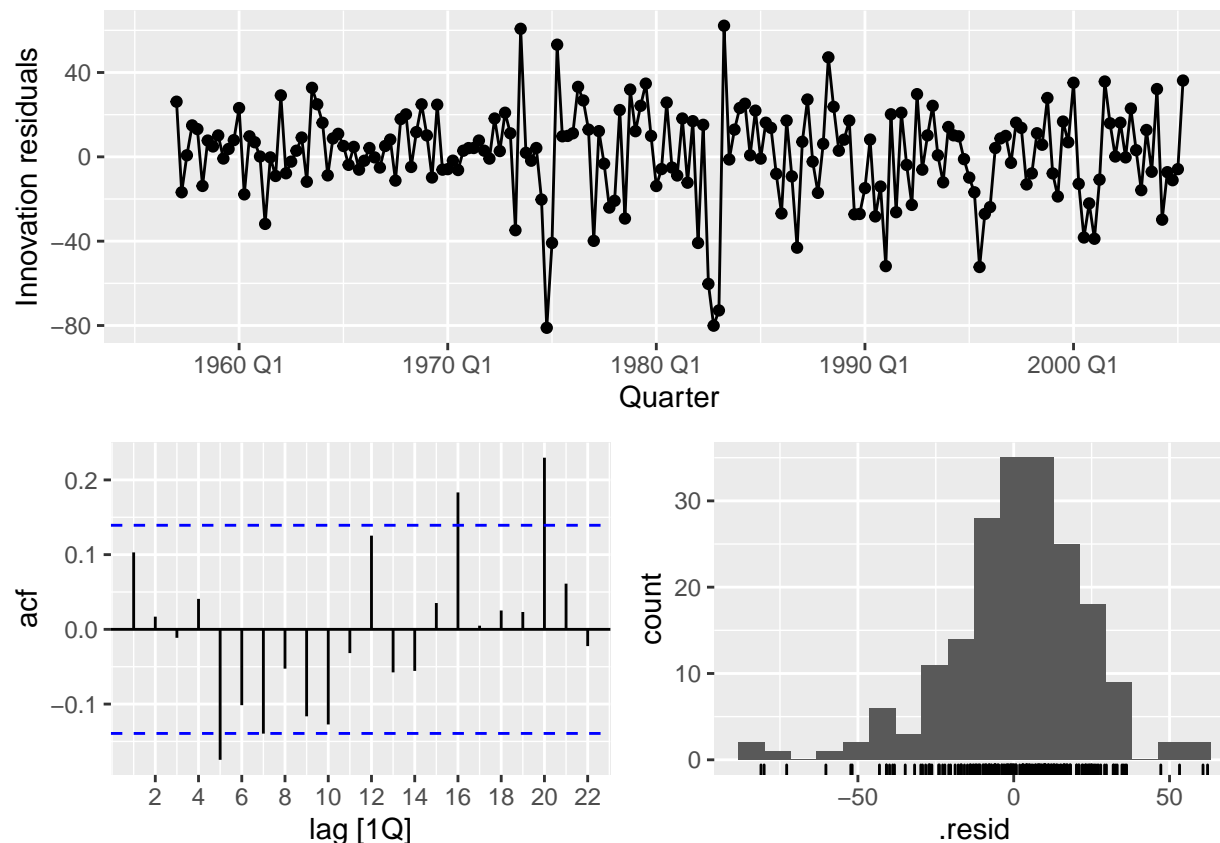


```
fit_dcmp %>% gg_tsresiduals()
```

```
## Warning: Removed 4 row(s) containing missing values (geom_path).
```

```
## Warning: Removed 4 rows containing missing values (geom_point).
```

```
## Warning: Removed 4 rows containing non-finite values (stat_bin).
```



There looks like there isn't much difference, But we can see that the residuals are less spread in the Robust STL than in the normal one, which means that it has captured more information.

g) Compare forecasts from `decomposition_model()` with those from `SNAIVE()`, using a test set comprising the last 2 years of data. Which is better?


```
bricks_train <- bricks_aus_production %>%
  filter(year(Quarter) <= 2002)
bricks_train
```

```
## # A tibble: 188 x 2 [1Q]
##   Quarter Bricks
##   <qtr>   <dbl>
## 1 1956 Q1    189
## 2 1956 Q2    204
## 3 1956 Q3    208
## 4 1956 Q4    197
## 5 1957 Q1    187
## 6 1957 Q2    214
## 7 1957 Q3    227
## 8 1957 Q4    222
## 9 1958 Q1    199
##10 1958 Q2    229
## # ... with 178 more rows
```

```
bricks_fit <- bricks_train %>%
  model(
    `Seasonal naïve` = SNAIVE(Bricks),
    stlf = decomposition_model(
      STL(Bricks ~ season(window = "periodic"), robust = TRUE), #, robust = TRUE
      NAIVE(season_adjust))
  )

bricks_fc <- bricks_fit %>%
  forecast(h = 10)

bricks_fc %>%
  autoplot(
    bricks_aus_production,
    level = NULL
  ) +
  labs(title = "Forecasts for Australian Bricks") +
  guides(colour = guide_legend(title = "Forecast"))
```



```
accuracy(bricks_fc, bricks_aus_production)
```

```
## # A tibble: 2 x 10
##   .model      .type    ME  RMSE   MAE   MPE  MAPE  MASE  RMSSE   ACF1
##   <chr>      <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
```

```
## 1 Seasonal naïve Test 17.8 29.2 23.8 4.32 5.82 0.656 0.592 -0.186
## 2 stlf Test -4.94 15.7 12.6 -1.31 3.17 0.347 0.318 0.0292
```

It looks like the `decomposition_model` achieve better metrics than the `SNAIVE` method in all the metrics and in the graph

## 5.11 exercises 12

a) Extract data from the Gold Coast region using `filter()` and aggregate total overnight trips (sum over Purpose) using `summarise()`. Call this new dataset `gc_tourism`.

```
tourism
```

```
## # A tibble: 24,320 x 5 [1Q]
## # Key:      Region, State, Purpose [304]
##   Quarter Region State Purpose Trips
##   <qtr> <chr> <chr> <chr> <dbl>
## 1 1998 Q1 Adelaide South Australia Business 135.
## 2 1998 Q2 Adelaide South Australia Business 110.
## 3 1998 Q3 Adelaide South Australia Business 166.
## 4 1998 Q4 Adelaide South Australia Business 127.
## 5 1999 Q1 Adelaide South Australia Business 137.
## 6 1999 Q2 Adelaide South Australia Business 200.
## 7 1999 Q3 Adelaide South Australia Business 169.
## 8 1999 Q4 Adelaide South Australia Business 134.
## 9 2000 Q1 Adelaide South Australia Business 154.
## 10 2000 Q2 Adelaide South Australia Business 169.
## # ... with 24,310 more rows
```

```
gc_tourism <- tourism %>%
  filter(Region == "Gold Coast") %>%
  #group_by(Purpose) %>%
  summarise(TotalTrips = sum(Trips)) %>%
  #select(Quarter, Purpose, TotalTrips)
gc_tourism
```

```
## # A tibble: 80 x 2 [1Q]
##   Quarter TotalTrips
##   <qtr> <dbl>
## 1 1998 Q1 827.
## 2 1998 Q2 681.
## 3 1998 Q3 839.
## 4 1998 Q4 820.
## 5 1999 Q1 987.
## 6 1999 Q2 751.
## 7 1999 Q3 822.
## 8 1999 Q4 914.
## 9 2000 Q1 871.
## 10 2000 Q2 780.
## # ... with 70 more rows
```

b) Using `slice()` or `filter()`, create three training sets for this data excluding the last 1, 2 and 3 years. For example, `gc_train_1 <- gc_tourism %>% slice(1:(n()-4))`.

```
gc_train_1 <- gc_tourism %>% slice(1:(n()-4))
gc_train_1
```

```
## # A tsibble: 76 x 2 [1Q]
##   Quarter TotalTrips
##   <qtr>      <dbl>
## 1 1998 Q1      827.
## 2 1998 Q2      681.
## 3 1998 Q3      839.
## 4 1998 Q4      820.
## 5 1999 Q1      987.
## 6 1999 Q2      751.
## 7 1999 Q3      822.
## 8 1999 Q4      914.
## 9 2000 Q1      871.
##10 2000 Q2      780.
## # ... with 66 more rows
```

```
gc_train_2 <- gc_tourism %>% slice(1:(n()-8))
gc_train_2
```

```
## # A tsibble: 72 x 2 [1Q]
##   Quarter TotalTrips
##   <qtr>      <dbl>
## 1 1998 Q1      827.
## 2 1998 Q2      681.
## 3 1998 Q3      839.
## 4 1998 Q4      820.
## 5 1999 Q1      987.
## 6 1999 Q2      751.
## 7 1999 Q3      822.
## 8 1999 Q4      914.
## 9 2000 Q1      871.
##10 2000 Q2      780.
## # ... with 62 more rows
```

```
gc_train_3 <- gc_tourism %>% slice(1:(n()-12))
gc_train_3
```

```
## # A tsibble: 68 x 2 [1Q]
##   Quarter TotalTrips
##   <qtr>      <dbl>
## 1 1998 Q1      827.
## 2 1998 Q2      681.
## 3 1998 Q3      839.
## 4 1998 Q4      820.
## 5 1999 Q1      987.
## 6 1999 Q2      751.
## 7 1999 Q3      822.
```

```
## 8 1999 Q4      914.
## 9 2000 Q1      871.
## 10 2000 Q2     780.
## # ... with 58 more rows
```

```
gc_test_1 <- gc_tourism %>% slice((n()-4+1):(n()))
gc_test_1
```

```
## # A tsibble: 4 x 2 [1Q]
##   Quarter TotalTrips
##   <qtr>      <dbl>
## 1 2017 Q1     1140.
## 2 2017 Q2      904.
## 3 2017 Q3     1053.
## 4 2017 Q4      908.
```

```
gc_test_2 <- gc_tourism %>% slice((n()-8+1):(n()-4))
gc_test_2
```

```
## # A tsibble: 4 x 2 [1Q]
##   Quarter TotalTrips
##   <qtr>      <dbl>
## 1 2016 Q1      933.
## 2 2016 Q2      854.
## 3 2016 Q3      850.
## 4 2016 Q4     1066.
```

```
gc_test_3 <- gc_tourism %>% slice((n()-12+1):(n()-8))
gc_test_3
```

```
## # A tsibble: 4 x 2 [1Q]
##   Quarter TotalTrips
##   <qtr>      <dbl>
## 1 2015 Q1      943.
## 2 2015 Q2      800.
## 3 2015 Q3      895.
## 4 2015 Q4     1017.
```

c) Compute one year of forecasts for each training set using the seasonal naïve (SNAIVE()) method. Call these `gc_fc_1`, `gc_fc_2` and `gc_fc_3`, respectively.

```
gc_fit_1 <- gc_train_1 %>%
  model(SNAIVE(TotalTrips))
# Generate forecasts for 4 quarters
gc_fc_1 <- gc_fit_1 %>% forecast(h = 4)
```

```
gc_fit_2 <- gc_train_2 %>%
  model(SNAIVE(TotalTrips))
# Generate forecasts for 4 quarters
gc_fc_2 <- gc_fit_2 %>% forecast(h = 4)
```

```
gc_fit_3 <- gc_train_3 %>%
  model(SNAIVE(TotalTrips))
# Generate forecasts for 4 quarters
gc_fc_3 <- gc_fit_3 %>% forecast(h = 4)
```

d) Use `accuracy()` to compare the test set forecast accuracy using MAPE. Comment on these.

```
accuracy(gc_fc_1, gc_test_1)
```

```
## # A tibble: 1 x 10
##   .model      .type    ME  RMSE  MAE  MPE  MAPE  MASE  RMSSE  ACF1
##   <chr>      <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 SNAIVE(TotalTrips) Test   75.1  167.  154.   6.36  15.1   NaN   NaN  -0.410
```

```
accuracy(gc_fc_2, gc_test_2)
```

```
## # A tibble: 1 x 10
##   .model      .type    ME  RMSE  MAE  MPE  MAPE  MASE  RMSSE  ACF1
##   <chr>      <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 SNAIVE(TotalTrips) Test   12.0  43.1  39.5   1.14   4.32   NaN   NaN  -0.792
```

```
accuracy(gc_fc_3, gc_test_3)
```

```
## # A tibble: 1 x 10
##   .model      .type    ME  RMSE  MAE  MPE  MAPE  MASE  RMSSE  ACF1
##   <chr>      <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 SNAIVE(TotalTrips) Test   35.8  91.4  83.9   3.56   9.07   NaN   NaN   0.239
```

If we take a look at the accuracy using MAPE, we can see that for the model 2 was the lowest value at 4.320729, follow by the model 3 at 9.067368, while model 1 got the highest value at 15.06055.